

BIG DATA ANALYSIS USING DASK

Project Title

Big Data Analysis on Social Network Ads Dataset using Dask

Objective

To perform scalable analysis on a dataset using Dask and extract insights such as:

- Total rows and columns
- Missing values
- Purchased distribution
- Gender distribution
- Group-wise average salary and age
- Top purchased users

Tool Used

Dask DataFrame

- Dask supports parallel processing using partitions.
- It is suitable for scalable big data processing.

```
In [8]: import dask.dataframe as dd

# Load dataset
df = dd.read_csv("social.csv")    # keep your file name correct
df.compute()
```

Out[8]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1
8	15600575	Male	25	33000	0
9	15727311	Female	35	65000	0
10	15570769	Female	26	80000	0
11	15606274	Female	26	52000	0
12	15746139	Male	20	86000	0
13	15704987	Male	32	18000	0
14	15628972	Male	18	82000	0
15	15697686	Male	29	80000	0
16	15733883	Male	47	25000	1
17	15617482	Male	45	26000	1
18	15704583	Male	46	28000	1
19	15621083	Female	48	29000	1
20	15649487	Male	45	22000	1
21	15736760	Female	47	49000	1
22	15714658	Male	48	41000	1
23	15599081	Female	45	22000	1
24	15705113	Male	46	23000	1
25	15631159	Male	47	20000	1
26	15792818	Male	49	28000	1
27	15633531	Female	47	30000	1
28	15744529	Male	29	43000	0
29	15669656	Male	31	18000	0
30	15581198	Male	31	74000	0
31	15729054	Female	27	137000	1
32	15573452	Female	21	16000	0
33	15776733	Female	28	44000	0
34	15724858	Male	27	90000	0
35	15713144	Male	35	27000	0

	User ID	Gender	Age	EstimatedSalary	Purchased
36	15690188	Female	33	28000	0
37	15689425	Male	30	49000	0
38	15671766	Female	26	72000	0
39	15782806	Female	27	31000	0
40	15764419	Female	27	17000	0
41	15591915	Female	33	51000	0
42	15772798	Male	35	108000	0
43	15792008	Male	30	15000	0
44	15715541	Female	28	84000	0
45	15639277	Male	23	20000	0
46	15798850	Male	25	79000	0
47	15776348	Female	27	54000	0
48	15727696	Male	30	135000	1
49	15793813	Female	31	89000	0
50	15694395	Female	24	32000	0
51	15764195	Female	18	44000	0

```
In [13]: print("Total Rows:", df.shape[0].compute())
print("Total Columns:", len(df.columns))
print("Partitions (Parallelism):", df.npartitions)

print(df.dtypes)
```

```
Total Rows: 52
Total Columns: 5
Partitions (Parallelism): 1
User ID          int64
Gender           string
Age              int64
EstimatedSalary  int64
Purchased        int64
dtype: object
```

Missing Values Analysis

We check how many missing values are present in each column.

```
In [10]: missing = df.isnull().sum().compute()
print("Missing Values:\n", missing)
```

```
Missing Values:
User ID          0
Gender           0
Age              0
EstimatedSalary  0
Purchased        0
dtype: int64
```

Purchased Distribution

This shows how many users purchased (1) and not purchased (0).

In [12]: `df[df["Purchased"] == 1].compute()`

Out[12]:

	User ID	Gender	Age	EstimatedSalary	Purchased
7	15694829	Female	32	150000	1
16	15733883	Male	47	25000	1
17	15617482	Male	45	26000	1
18	15704583	Male	46	28000	1
19	15621083	Female	48	29000	1
20	15649487	Male	45	22000	1
21	15736760	Female	47	49000	1
22	15714658	Male	48	41000	1
23	15599081	Female	45	22000	1
24	15705113	Male	46	23000	1
25	15631159	Male	47	20000	1
26	15792818	Male	49	28000	1
27	15633531	Female	47	30000	1
31	15729054	Female	27	137000	1
48	15727696	Male	30	135000	1

In [14]: `df[df["Purchased"] == 0].compute()`

Out[14]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
8	15600575	Male	25	33000	0
9	15727311	Female	35	65000	0
10	15570769	Female	26	80000	0
11	15606274	Female	26	52000	0
12	15746139	Male	20	86000	0
13	15704987	Male	32	18000	0
14	15628972	Male	18	82000	0
15	15697686	Male	29	80000	0
28	15744529	Male	29	43000	0
29	15669656	Male	31	18000	0
30	15581198	Male	31	74000	0
32	15573452	Female	21	16000	0
33	15776733	Female	28	44000	0
34	15724858	Male	27	90000	0
35	15713144	Male	35	27000	0
36	15690188	Female	33	28000	0
37	15689425	Male	30	49000	0
38	15671766	Female	26	72000	0
39	15782806	Female	27	31000	0
40	15764419	Female	27	17000	0
41	15591915	Female	33	51000	0
42	15772798	Male	35	108000	0
43	15792008	Male	30	15000	0
44	15715541	Female	28	84000	0
45	15639277	Male	23	20000	0
46	15798850	Male	25	79000	0
47	15776348	Female	27	54000	0
49	15793813	Female	31	89000	0
50	15694395	Female	24	32000	0

User ID	Gender	Age	EstimatedSalary	Purchased
51	15764195	Female	18	44000

Gender Distribution

This shows Male vs Female count.

```
In [15]: gender_count = df["Gender"].value_counts().compute()
print("Gender Distribution:\n", gender_count)
```

```
Gender Distribution:
Gender
Female    24
Male     28
Name: count, dtype: int64[pyarrow]
```

Group-wise Analysis

We find average Age and Estimated Salary for Purchased=0 and Purchased=1.

```
In [16]: grouped = df.groupby("Purchased")[["Age", "EstimatedSalary"]].mean().compute()
print("Average Age & Salary by Purchased:\n", grouped)
```

```
Average Age & Salary by Purchased:
          Age  EstimatedSalary
Purchased
0           27.297297      52378.378378
1           43.266667      51000.000000
```

Top Purchased Users

We extract top 5 users with highest salary among Purchased = 1.

```
In [17]: top_salary = df[df["Purchased"] == 1][["User ID", "Gender", "Age", "EstimatedSalary"]] \
    .nlargest(5, "EstimatedSalary") \
    .compute()

print("Top 5 highest salary purchased users:\n")
print(top_salary)
```

Top 5 highest salary purchased users:

	User ID	Gender	Age	EstimatedSalary
7	15694829	Female	32	150000
31	15729054	Female	27	137000
48	15727696	Male	30	135000
21	15736760	Female	47	49000
22	15714658	Male	48	41000

Final Insights & Conclusion

- Majority of users did not purchase.
- Purchased users generally have higher salary and higher age.
- Dask processed the dataset efficiently using partitions, showing scalability for big data.

```
In [18]: print("----- FINAL INSIGHTS -----")
print("1) Purchased Distribution:\n", purchase_count)
```

```
print("\n2) Gender Distribution:\n", gender_count)
print("\n3) Avg Age & Salary by Purchased:\n", grouped)
print("\nScalability proof - Partitions:", df.npartitions)
```

----- FINAL INSIGHTS -----

1) Purchased Distribution:

Purchased

0 37

1 15

Name: count, dtype: int64

2) Gender Distribution:

Gender

Female 24

Male 28

Name: count, dtype: int64[pyarrow]

3) Avg Age & Salary by Purchased:

Age EstimatedSalary

Purchased

0 27.297297 52378.378378

1 43.266667 51000.000000

Scalability proof - Partitions: 1